Few-shot Learning	Few-shot learning refers to the training of machine learning algorithms using a very small set of training data instead of a very large set. This is most suitable in the field of computer vision, where it is desirable to have an object categorization model work well without thousands of training examples.
Flume	Flume is a service designed for streaming logs into the Hadoop environment. It can collect and aggregate huge amounts of log data from a variety of sources. In order to collect high volume of data, multiple flume agents can be configured. Here are the major features of Apache Flume: • Flume is a flexible tool as it allows to scale in environments with as low as five machines to as high as several thousands of machines • Apache Flume provides high throughput and low latency • Apache Flume has a declarative configuration but provides ease of extensibility • Flume in Hadoop is fault tolerant, linearly scalable and stream oriented
Frequentist Statistics	**Frequentist Statistics** tests whether an event (hypothesis) occurs or not. It calculates the probability of an event in the long run of the experiment (i.e the experiment is repeated under the same conditions to obtain the outcome). Here, the sampling distributions of **fixed size** are taken. Then, the experiment is theoretically repeated **infinite number of times** but practically done with a stopping intention. For example, I perform an experiment with a stopping intention in mind that I will stop the experiment when it is repeated 1000 times or I see minimum 300 heads in a coin toss....